

Skill registry — 2026-05-20

Scope: every `.claude/skills/*/SKILL.md` across the portfolio. Rows tagged **measured** show real token consumption from Claude Code transcripts (90-day window), with the annual figure projected linearly. Rows tagged **estimated** are author guesses parsed from each repo's docs/README. On measured rows, a third line compares the observed average to the author's prior estimate, e.g. `est. 4K • 93x under`.

Reference: Claude Code built-ins

These are Claude Code's own built-in slash commands — **not part of this portfolio**. Shown here purely so the custom-skill numbers below have a familiar scale to compare against. **Excluded from every total on the page.**

Skill	Description	Tokens / use	Uses	Total	Receipt
<code>/review</code>	Claude Code built-in PR code review	1.10M (observed)	13 in 90d → ~53/yr proj.	observed 14.26M proj. ~58.14M/yr	measured last 2026-05-20

Aggregate

Repo	Skills	Redirects	With receipts	Observed (90d)	Tokens / yr (proj.)
AudiobookMaker	8	0	6	~202K	~2.72M
Spacepotatis	12	1	11	~930K	~5.06M
claude-skills	9	0	9	~2.62M	~21.67M
mikkonumminen.dev	4	0	3	~201K	~1.06M
Total	33	1	32	~3.95M	~30.51M

AudiobookMaker — <https://github.com/MikkoNumminen/AudiobookMaker>

Skill	Description	Tokens / use	Uses	Total	Receipt
ci-failure-triage	When CI fails on a release tag or PR, walk the failure modes in the right order against a recipe library built from the project's <code>fix(ci)</code> commit history. Use whenever the user says "CI is red", "the build failed", "tag failed CI", "release-build broke", or any time gh run list shows a recent failure on master or a release tag.	1K (est.)	88 / yr (est.)	~88K (est.)	estimated

Skill	Description	Tokens / use	Uses	Total	Receipt
copyright-scan	Scan a git diff (staged by default, or a named revision range / file set) for accidental third-party copyright leaks before they land on origin. Use this skill whenever the user says "scan the diff", "copyright check", "scan for leaks", "is this safe to push", "check this before I commit", "did I leak anything?", or whenever you are about to create a commit or push in this repo. CLAUDE.md marks copyright leakage as a P0 — this skill turns the manual scan ritual into one pass. Produces a structured pass/fail report with file:line citations and, on fail, the exact remediation path (edit vs. un-stage vs. P0 scrub-from-origin).	14K (observed) est. 3K · 4.6× under	1 in 90d → ~4/yr proj.	observed 14K proj. ~56K/yr	measured last 2026-05-12
pronunciation-corpus-add	Append a Finnish pronunciation failure (a word the Chatterbox Grandmom voice mispronounces) to the project's pronunciation corpus at docs/pronunciation_corpus.fi.md. Use this whenever the user or a tester (Turo is the canonical one) reports that Grandmom said a Finnish word wrong — e.g. "Grandmom pronounces X as Y", "add this to the corpus", "log this pronunciation bug", "the TTS messed up <word>". Also use when the user pastes a transcript of what was heard vs. what was written. This corpus is the evidence base for targeted Pass I / lexicon fixes; entries must be structured so patterns across reports are visible at a glance.	—	0 / yr (est.)	—	estimated
release-bundle-audit	Audit the AudiobookMaker PyInstaller release bundle for unused dependencies, dead-code data files, and accidental ML-stack pollution; then propose spec-only fixes on a `chore/release-bundle-size` branch. Use whenever the user says "the installer is huge", "why is the bundle so big", "shake out unused deps before release", "shrink the .exe", "audit the .spec", or asks why a frozen build ships some specific package. Encodes the empirical finding that ~28% of the v3.11.x bundle was dead-code data files (ffplay.exe, AVIF support, Arabic phonemizer, ONNX training tools) plus a defense-in-depth exclude set that prevents future regressions when a contributor accidentally pip-installs torch on the build machine.	3K (est.)	4 / yr (est.)	~12K (est.)	estimated
release-cut	Cut a new AudiobookMaker release (bump APP_VERSION, tag vX.Y.Z, push, verify CI lands SHA-256 in release notes AND sidecar). Use this skill whenever the user says "cut a release", "bump the version", "ship X.Y.Z", "make a release", "release X.Y.Z", or asks to tag a new version. Auto-update is existential for this project — a release that ships without a valid SHA-256 locks every existing user out of one-click updates. This skill encodes every moving part so that can never happen again.	188K (observed)	1 in 90d → ~4/yr proj.	observed 188K proj. ~753K/yr	measured last 2026-05-12
voice-pack-finnish	Build a Finnish voice pack from a source audio file end-to-end — chunked analyze with ECAPA diarization, transcript validation per speaker, branching to LoRA training (long sources) or few-shot ref-clip packaging (short sources), and ear-check by synth. Use whenever the user says "copy the voices", "extract voices from this clip", "make a voice pack from this audio", "I have a short Finnish clip", or hands you a Finnish podcast / interview / audiobook to mimic. Encodes the empirically-validated Finnish pipeline. CRITICAL — Finnish source only (ASR + finetune model are Finnish-specific); for English / other languages this skill does not apply. Never assume single-speaker, never trust pyannotate labels without validation, never install copyright-derived packs to ~/.audiobookmaker.	7K (est.)	248 / yr (est.)	~1.61M (est.)	estimated

Skill	Description	Tokens / use	Uses	Total	Receipt
work-session	Start, pause, or finish a TODO.md work session as one of the 4 permanent Claude sessions in AudiobookMaker. Use this whenever the user says "claim task X", "take task Y", "start working on Z", "pick something", "I'm done", "finish up", "pause", "I'm blocked", or "go idle". Also use before touching any code on a fresh session to make sure the status board and in-progress list are honest. Parallel Claudes collide without this discipline; the shared TODO.md protocol is the only way four sessions stay coherent.	2K (est.)	—	—	estimated
worktree-launch	Start a new parallel Claude session safely. Pick a free session slot (Claude 1/2/3/4), create the worktree under .claude/worktrees/<branch>, claim a task in TODO.md, verify isolation actually held after the agent runs. Use whenever the user says "start a new session", "spawn another Claude", "open a worktree", or "let me run two of you in parallel".	800 (est.)	252 / yr (est.)	~202K (est.)	estimated

Spacepotatis — <https://github.com/MikkoNumminen/Spacepotatis>

Skill	Description	Tokens / use	Uses	Total	Receipt
balance-review	Diff uncommitted changes to game data (weapons, enemies, waves, missions, perks, augments, loot pools, solar systems) and report DPS, TTK, energy-cost-per-DPS, augment-folded effective DPS, loot-pool roster shifts, and drop-rate deltas vs HEAD.	14K (est.)	50 / yr (est.)	~675K (est.)	estimated
content-audit	Pre-commit content invariants check — orphan refs (enemy / weapon / sprite / pod / loot-pool / mission system / story trigger), missing sprite generators, perk drop-weight sanity, mission prereq DAG, story integrity (voice + music files, trigger refs), storyTriggers helper coverage.	53K (observed) est. 15K · 3.5× under	1 in 90d → ~4/yr proj.	observed 53K proj. ~213K/yr	measured last 2026-05-03
equipment	Create, modify, or remove a weapon or piece of equipment (augment / reactor / shield / armor) — including visual changes to how it appears in combat or in the loadout UI. Covers stats, sprites, prices, and the full CRUD lifecycle for the entire ship-loadout content surface.	401K (observed) est. 4K · 93× under	2 in 90d → ~8/yr proj.	observed 803K proj. ~3.21M/yr	measured last 2026-05-03
modular-architecture-audit	Orchestrates the multi-phase modular-architecture audit + refactor. Phase 0 sets up agents; Phase 1 inventories every file; Phase 2 proposes module boundaries; Phase 3 mechanically extracts modules one at a time (parallelizable across worktrees); Phase 4 writes AI-first documentation; Phase 5 verifies. Each phase produces a written artifact under docs/audit/ and STOPS at a gate for user approval. Never auto-advances.	5K (est.)	5 / yr (est.)	~25K (est.)	estimated
new-enemy	Scaffold a new enemy — enemies.json entry, BootScene placeholder sprite generator, optional test wave, and integrity test verification.	6K (est.)	25 / yr (est.)	~138K (est.)	estimated
new-migration	Add a Postgres schema migration end-to-end — dated SQL file in db/migrations/, Database interface update in src/lib/db.ts, prod application via scripts/migrate.mjs, schema verification via scripts/check-schema.mjs, and the PR-body checkbox that gates merge. Enforces CLAUDE.md \$7a HARD RULE so the next save POST doesn't 500.	7K (est.)	15 / yr (est.)	~105K (est.)	estimated

Skill	Description	Tokens / use	Uses	Total	Receipt
new-mission	Scaffold a new combat mission across missions.json, waves.json, the galaxy planet binding, and a smoke test — in one shot.	8K (est.)	30 / yr (est.)	~240K (est.)	estimated
new-perk	Scaffold a new mission-only perk — perks.ts entry, BootScene icon generator, HUD chip wiring, and (for actives) a switch case in PerkController.apply / triggerActive.	9K (est.)	10 / yr (est.)	~90K (est.)	estimated
new-solar-system	Add a new solar system to the galaxy — solarSystems.json entry (incl. galaxy bed), SolarSystemId union extension, optional unlock gating, REQUIRED on-system-enter cinematic (voice + music) AND dedicated galaxy-view music bed, and mission-binding TODO list.	13K (est.)	5 / yr (est.)	~65K (est.)	estimated
new-story	Add, modify, or remove story content — cinematic popups, voiceovers, music beds, body text, written synopses, and auto-trigger wiring. Covers the full CRUD lifecycle for the Story log.	56K (observed) est. 5K · 10× under	1 in 90d → ~4/yr proj.	observed 56K proj. ~222K/yr	measured last 2026-04-30
new-weapon <i>(redirect)</i>	Superseded by /equipment. Adding a weapon now lives inside the broader CRUD-and-visuals skill. This stub redirects.	—	—	—	—
save-roundtrip-audit	Pre-commit save-pipeline invariants check — walks every StateSnapshot field through the 8 layers of the save round-trip (snapshot interface -> toSnapshot -> POST schema -> POST handler -> DB column -> migration -> GET handler -> RemoteSaveSchema -> sync.loadSave -> hydrate) and flags any layer that silently drops the field. Read-only.	19K (observed) est. 12K · 1.5× under	1 in 90d → ~4/yr proj.	observed 19K proj. ~74K/yr	measured last 2026-05-03

claude-skills — <https://github.com/MikkoNumminen/claude-skills>

Skill	Description	Tokens / use	Uses	Total	Receipt
ai-codegen-smell-audit	Read-only audit for specific failure modes that recur in LLM-generated code — defensive guards on impossible cases, generic names in domain code, swallowed errors, single-use helpers, mirror tests, and so on. Each check has a concrete smell example and a concrete legitimate counter-example so the auditor (human + assistant) can tell signal from noise. Produces `docs/audits/ai-smell-<YYYY-MM-DD>.md` with severity-ranked findings.	199K (observed) est. 75K · 2.7× under	1 in 90d → ~4/yr proj.	observed 199K proj. ~798K/yr	measured last 2026-05-17
audit	Run a multi-phase robustness audit of the current codebase. Phase 1 tries language-appropriate static-analysis tools, skipping cleanly if absent. Phase 2 spawns five parallel subagents across non-overlapping scopes — resource lifecycle, data integrity, concurrency, error paths, external boundaries — each returning findings with file:line citations and severity. Phase 3 aggregates into docs/audits/audit-YYYY-MM-DD.md with a severity tally (critical / high / medium / low) and a branch topology recommendation for the fix follow-up.	553K (observed) est. 425K · 1.3× under	2 in 90d → ~8/yr proj.	observed 1.11M proj. ~4.42M/yr	measured last 2026-05-20

Skill	Description	Tokens / use	Uses	Total	Receipt
mikko-audit-suite	Run every `mikko-*` audit skill that's relevant to the current codebase, in a sensible order, and produce one index document linking to each audit's report. Detects the codebase shape (language, framework, security surface), maps to the suggested audit list, confirms once with the human, then runs each audit sequentially. Expensive — total tokens are the sum of every dispatched audit (typically ~100-500K).	300K (est.)	18 / yr (est.)	~5.40M (est.)	estimated
mikko-help	List every installed `mikko-*` skill with its description (or with its `barney:` field when `--barney` is passed). The fast answer to "I know I have a skill for this but can't remember which one". Pass `--detect` for codebase-aware audit recommendations. Pass `--barney` for plain-English short descriptions. No sub-agents, no network, no measurements — just a table.	644K (observed) est. 7K · 99× under	2 in 90d → ~8/yr proj.	observed 1.29M proj. ~5.15M/yr	measured last 2026-05-20
mikko-install	Install, update, list, or uninstall `mikko-*` skills from the cloned `claude-skills` source repo into the user-wide skill directory (`~/claude/skills/mikko-*`). Wraps the existing `install-mikko.sh` and `install.sh` scripts so the install loop is reachable from inside any Claude Code session via a slash command. Does NOT clone the repo. Does NOT touch skills outside the `mikko-*` prefix.	4K (est.)	5 / yr (est.)	~20K (est.)	estimated
react-anti-patterns-audit	Read-only audit for six concrete React anti-patterns that recur in component code (key=index in lists, useEffect for derived state, dep-array lies, state mutation instead of replace, missing effect cleanup, multiple sources of truth). Each check has a concrete smell example and a concrete legitimate counter-example. Runs a pre-flight check first — aborts cleanly if the codebase isn't React.	23K (est.)	30 / yr (est.)	~675K (est.)	estimated
readme-drift-sync	Detect drift between what the repo actually contains and what `README.md` claims, then rewrite ONLY the drifted sections in the README's existing voice. Five drift axes (file structure, dependencies, skills, features, status). Voice-preserving. Write-with-approval: emits a working-tree edit to `README.md` plus a drift report and a voice-profile cache, then STOPS. Does NOT commit, does NOT push.	40K (est.)	15 / yr (est.)	~600K (est.)	estimated
security-audit	Orchestrates the multi-phase security audit + remediation. Phase 0 sets up agents; Phase 1 maps the full attack surface (entry points, auth, authz, input, secrets, data exposure, deps, transport, client, ops); Phase 2 turns findings into a prioritized remediation plan; Phase 3 fixes one finding at a time with regression tests; Phase 4 writes AI-first security documentation; Phase 5 verifies. Each phase produces an artifact under docs/security/ and STOPS at a gate for user approval. Never auto-advances.	750K (est.)	6 / yr (est.)	~4.50M (est.)	estimated
skill-usage	Measure actual Claude Code skill usage from local transcript JSONL files. Walks `~/claude/projects/*/*.jsonl` (and their `subagents/*/*.jsonl` sidechain files), filters assistant messages by the harness-emitted `attributionSkill` field, groups invocations by `(sessionId, skill)`, sums tokens deduped by `requestId`. Emits a dated `SKILL-USAGE-{YYYY-MM-DD}.json` with per-skill measured invocation counts and token totals.	13K (observed) est. 4K · 3.3× under	2 in 90d → ~8/yr proj.	observed 27K proj. ~106K/yr	measured last 2026-05-20

Skill	Description	Tokens / use	Uses	Total	Receipt
md-to-pdf	Render a markdown report (or any HTML) to a styled PDF using the local Chrome's `--print-to-pdf` flag. Zero npm install — no puppeteer / Chromium download. Claude assembles a tuned HTML file in-context (markdown → HTML conversion + per-document CSS), then invokes `scripts/build-pdf.mjs` to land the PDF. Used by the contact-page terminal's `download --skills` (registry PDF) and any one-off "I need this report as a PDF" moment.	25K (est.)	10 / yr (est.)	~250K (est.)	estimated
skill-pdf	One-command refresh of the skills-registry PDF. Wraps a pre-flight check plus four sequenced actions that take a recent token-measurement snapshot and produce a deploy-ready `public/skills-registry.pdf`: (1) `/skill-registry` to re-walk the portfolio and emit a fresh inventory JSON, (2) `npm run sync:skills-registry` to copy the LATEST inventory into `public/data/`, (3) `node scripts/apply-measurement-overlay.mjs` to merge transcript measurements into the inventory and dedupe library-canonical duplicates, (4) `npm run build:skills-pdf` to render the HTML and print to PDF via local Chrome. Use whenever the user says "refresh the skills PDF", "rebuild the skill-registry PDF", "regenerate the skills-registry document", "publish the skills PDF", or `/skill-pdf`. Does NOT gather token usage — that's `/mikko-skill-usage`, run separately AND from inside `mikkonumminen.dev/` so the JSON lands at the path this skill expects. Does NOT commit or push.	83K (est.)	—	—	estimated
skill-registry	Scan every sibling repo under D:/koodaamista for `.claude/skills/*/SKILL.md` files and emit a consolidated registry as a structured JSON document — per-skill name, description, redirect flag, and (where receipts exist) token-savings estimates. Runs one Sonnet sub-agent per repo in parallel for swiftness. Output goes to a dated JSON file under `.claude/agent-verdicts/`, committed so other Claude sessions can read the current inventory without re-running.	83K (observed) est. 95K · 1.1× over	1 in 90d → ~4/yr proj.	observed 83K proj. ~331K/yr	measured last 2026-05-19
sync-readmes	Audit project data against sibling repos' READMEs and open a PR with drift corrections. Runs parallel Sonnet diff agents (one per sibling repo), synthesizes drift, applies en+fi+sv corrections plus tech-list updates, and lands a PR for review on GitHub.	119K (observed) est. 141K · 1.2× over	1 in 90d → ~4/yr proj.	observed 119K proj. ~475K/yr	measured last 2026-05-18