

Skill-suite calibration — mikko- library (2026-06-02) + Spacepotatis audit skills (2026-06-03)

green = saved ≥10% · red = cost ≥10% more · black = within ±10% (no signed signal). Colours the “% saved” columns only; swing/pp and token columns stay neutral.

A/B token measurement (cold arm A vs with-skill arm B) for the 8 cleanly-A/B-able mikko- skills, across **all three models** (Opus 4.8, Sonnet 4.6, Haiku 4.5). This is the “before” baseline — the skills were just audited + given freshness blocks (claude-skills #22/#23); a planned optimization pass will re-measure against these numbers. The optimization pass (claude-skills #24) has since landed and three of these skills were re-measured — see **After-optimization re-measure** below. Method mirrors the first calibration document, extended to the full updated suite and every model.

Method

- **Two arms, same task.** Arm A (cold) scouts from first principles; arm B reads the `SKILL.md`, runs any companion script, follows the procedure. Same deliverable per pair. **48 sub-agent arms total** (8 skills × 2 arms × 3 models).
- **Corpus (8):** `audit`, `ai-codegen-smell-audit`, `react-anti-patterns-audit`, `readme-drift-sync`, `skills-quality`, `skills-freshness`, `skill-usage`, `session-cost`. (Orchestrator/installer/lister/recursive skills excluded as not cleanly measurable.)
- **Tasks auto-synthesized**, pinned to fixed targets: code audits → `mikkonumminen.dev/src/lib` (+ `/terminal`); react → `Spacepotatis`; readme → `mikkonumminen.dev/README.md`; skills-quality/freshness → live `~/ .claude/skills/mikko-*`; skill-usage/session-cost → `~/ .claude/projects`.
- **Single-agent approximation** for orchestrators (`audit`, `readme-drift-sync` run their parallel passes inline). **Accounting:** per-arm `subagent_tokens` (input+output+cache-creation). Read-only. N = 1 per cell.

Results — all three models

Skill	Opus A	Opus B	Opus %	Sonnet A	Sonnet B	Sonnet %	Haiku A	Haiku B	Haiku %
skills-freshness	101,969	23,215	+77%	82,177	18,083	+78%	88,412	29,948	+66%
skills-quality	75,323	21,829	+71%	83,680	14,895	+82%	55,753	23,637	+58%
session-cost	31,926	22,572	+29%	19,823	15,859	+20%	69,520	23,054	+67%
react-anti-patterns-audit	103,254	97,460	+6%	73,050	68,354	+6%	55,856	57,445	−3%
audit	94,777	142,666	−51%	130,314	127,089	+2%	64,032	59,455	+7%
skill-usage	22,946	32,014	−40%	21,899	21,938	~0%	28,879	28,147	+3%
readme-drift-sync	33,383	44,341	−33%	24,539	74,709	−204%	41,126	39,715	+3%
ai-codegen-smell-audit	34,224	48,665	−42%	26,908	52,222	−94%	31,588	54,926	−74%*

Skill	Opus A	Opus B	Opus %	Sonnet A	Sonnet B	Sonnet %	Haiku A	Haiku B	Haiku %
Aggregate (ratio-of-sums)	497,802	432,762	+13%	462,390	393,149	+15%	435,166	316,327	+27%

* `ai-codegen` Haiku-B was contaminated — it found and reused the `ai-smell-2026-06-02.md` report a *Sonnet* arm wrote earlier (a side-effect-file bleed). Treat that cell as unreliable.

What the data shows (cross-model)

- Savings scale inversely with model capability** — aggregate +13% (Opus) → +15% (Sonnet) → +27% (Haiku). A weaker model scouts cold less efficiently, so the skill's deterministic procedure / script wins more, relatively. This is the optimization study's and first calibration's core thesis, reproduced cleanly across the whole suite.
- audit is the textbook monotonic case** — −51% → +2% → +7%. On Opus a cold scout is cheap (95K), so the orchestrator-run skill (143K) loses badly; on Haiku the cold scout is dearer relative to the skill, so it flips positive.
- The savings live in the three script-backed skills** — `skills-quality` (+71/+82/+58%), `skills-freshness` (+77/+78/+66%), `session-cost` (+29/+20/+67%) save on every model. They replace LLM reasoning with a Python pre-pass / `scan.mjs`. Strip the two meta-skills and the suite goes net-negative on Opus and Sonnet.
- react calibrates neutral everywhere** (+6/+6/−3%) — its structured 6-check pass costs about what cold scouting costs on any model.
- The big meta-skill saves are "save by not looking."** On all three models the skill arms short-circuited (pre-pass / sha256 hash → "nothing to review") while the **cold arms caught real issues**: bloat (`ai-codegen` 655 lines, `audit` 's 5× duplicated prompt boilerplate), `security-audit` targeting a non-existent codebase, broken cross-repo links, a stale "13 skills" count. The savings are partly a coarser read.
- Haiku is the least reliable arm** — it hallucinated a finding (`skill-calibration` "has no freshness check" — it does), wrongly concluded "token data unavailable" on skill-usage, and one Haiku arm tripped a security flag (ran PowerShell through Bash, circumventing a deny rule). Its verdicts carry the least confidence — same caveat the first calibration flagged.

Headline: across all three models the library's token economy is carried by the **script-backed skills** plus the neutral `react`; the prose audits are wash-to-negative against a capable cold model and only turn positive as the model weakens. The cold-arm findings (the bloat list, the real drift) are the targets for the upcoming optimization pass — which should re-measure against these numbers.

After-optimization re-measure

The "after" the baseline anticipated. `claude-skills` #24 optimized three of these skills — `ai-codegen-smell-audit` (655→580 lines; Provenance extracted to a companion file + 5 dead links removed), `audit` (410→390; the 5×-duplicated sub-agent prompt footer deduped), `readme-drift-sync` (325→310; dated content-calibration narrative extracted). Each optimized skill's **with-skill arm (B)** was re-measured on all three models — same tasks, same accounting. This compares **B(before-optim)** → **B(after-optim)**; a negative Δ means cheaper after.

Skill (arm B)	Opus →	Δ	Sonnet →	Δ	Haiku →	Δ
ai-codegen-smell-audit	48,665 → 50,044	+3%	52,222 → 39,148	−25%	54,926 → 47,403	−14%

Skill (arm B)	Opus →	Δ	Sonnet →	Δ	Haiku →	Δ
audit	142,666 → 82,841	−42%	127,089 → 142,847	+12%	59,455 → 56,378	−5%
readme-drift-sync	44,341 → 43,303	−2%	74,709 → 37,044	−50%	39,715 → 40,704	+2%
Aggregate (3 skills)	235,672 → 176,188	−25%	254,020 → 219,039	−14%	154,096 → 144,485	−6%

Reading: the optimization's per-invocation token effect is below the N = 1 noise floor — this delta is not the optimization. The per-cell swings span −50% to +12% with no relationship to the 1–3 KB of `SKILL.md` body each skill shed. The same `audit` task cost 56K, 59K, 83K, 127K and 143K across these runs; that ±40–60K task-work variance dwarfs the few-KB body trim by 20–60×. The overall −16% (all nine cells, 643,788 → 539,712) is the *before* pass's high outliers — `audit` Opus 143K, `readme` Sonnet 75K — regressing toward the mean, not a measured saving.

What this means for skill optimization. Trimming `SKILL.md` body size buys ≈ 0 measurable per-invocation tokens for task-heavy skills: the one-time body read is a rounding error against the audit work itself. The payoff of the #24 optimizations is real but lives where this metric can't see it — the always-loaded `description` (charged on every matching turn, not just on invocation), context cleanliness, and correctness (the dead links / dead game-refs / stale narrative removed). To *measure* a body-size win you'd need a skill whose body dominates its task (none of these three) or N >> 1 to drop the noise floor below the signal.

Spacepotatis audit-skills calibration (2026-06-03)

A second corpus, same A/B method, run after the Spacepotatis skills were finetuned (quality + freshness, Spacepotatis #286). The 4 cleanly-A/B-able read-only audit skills were measured cold (A) vs with the **finetuned** skill (B) across all three models — `balance-review`, `content-audit`, `ai-codegen-smell-audit`, `save-roundtrip-audit`. The 8 generative scaffolders (`new-*`, `equipment`) + 2 orchestrators (`modular-architecture-audit`, `security-audit`) were finetuned too but **excluded from A/B** — they mutate the repo, so they aren't cleanly read-only measurable (same exclusion logic the mikko- suite used). 24 sub-agent arms (4 × 2 × 3). Same accounting (`subagent_tokens`), N = 1, pinned read-only targets. The B arms read the finetuned `SKILL.md` from a worktree off the merged master; the cold A arms are skill-independent and unchanged.

Skill (arm B reads the finetuned <code>SKILL.md</code>)	Opus A	Opus B	Opus %	Sonnet A	Sonnet B	Sonnet %	Haiku A	Haiku B	Haiku %
balance-review	81,911	46,403	+43%	60,719	25,715	+58%	61,852	37,892	+39%
content-audit	121,173	66,732	+45%	99,250	71,408	+28%	62,300	57,438	+8%
ai-codegen-smell-audit	103,964	126,094	−21%	85,301	93,056	−9%	61,452	62,478	−2%
save-roundtrip-audit	64,695	72,437	−12%	44,787	58,010	−30%	73,877	62,038	+16%
Aggregate (ratio-of-sums)	371,743	311,666	+16%	290,057	248,189	+14%	259,481	219,846	+15%

What this corpus shows:

1. **balance-review** is the standout (+43/+58/+39%) — its structured metric procedure (DPS / TTK / energy-per-DPS formulas + a flagged-issues checklist) is far cheaper than scouting the balance maths cold. It's the Spacepotatis analogue of the mikko- script-backed skills.

2. **ai-codegen** is a wash-to-negative (-21/-9/-2%) — the same prose-audit pattern as its mikko- copy, reproduced on a different repo. The with-skill arm also follows the skill's "write a report" step, adding output tokens.
3. **content-audit** saves (+45/+28/+8%); **save-roundtrip** is mixed (-12/-30/+16%) — content's checklist beats cold scouting on every model; save-roundtrip's layer-by-layer procedure costs *more* on Opus/Sonnet (which scout the save pipeline efficiently cold) but saves on Haiku.
4. **Net is a flat ~+15% (+16/+14/+15%), not monotonic.** Unlike the mikko- suite's clean +13 → +15 → +27 inverse-capability curve, this corpus blends strong savers (balance, content) with washes (ai-codegen, save), so the per-model spread collapses to a flat band. The durable lesson reproduces across both repos: **token savings concentrate in procedure/script-backed skills; prose audits wash against a capable cold model.**

Noise-floor demonstration (unintended, but the cleanest in either corpus). This B column was re-measured. A first pass accidentally read the *pre-finetune* skills (the local checkout lagged the #286 merge), giving an aggregate of +6/+8/+7%; re-running the same arms against the **finetuned** skills — content differing by only a few KB + a freshness block — gave +16/+14/+15%, an ~8-point swing (e.g. **content** -Opus alone moved 100.9K → 66.7K). Two compounding noise sources explain it, neither a finetune effect: (a) the two skill versions are near-identical — a few KB of body + a freshness block is far too little to move 8 points; and (b) the cold A-arms were *not* re-run, so each cell now pairs a run-1 A against a run-2 B, layering cross-session drift on top of the per-cell **N = 1** variance. The swing is therefore noise, not signal — an accidental, direct confirmation of this doc's own thesis that run-to-run variance dwarfs SKILL.md content. *Which* skill saves vs washes is stable across both runs; only the magnitudes wobble within the noise band.

Caveats (this corpus): N = 1 per cell; the with-skill audit arms wrote stray `docs/audits/*-2026-06-03.md` reports during measurement (cleaned up); a Haiku **content-audit** arm tripped the PowerShell-via-Bash deny flag (read-only file listing); the generative + orchestrator skills were finetuned but not A/B-measured. Outcome ≠ tokens — e.g. one Haiku **save** arm flagged a "critical **currentPlanet** drop" that the Opus/Sonnet arms (and the skill) correctly read as a by-design re-derivation.

Caveats

- **N = 1 per cell** — direction + rough magnitude; prose-audit magnitudes are noisy (see `readme`: -33/-204/+3%).
- **Side-effect contamination** — skill arms wrote reports (`react-anti-patterns-2026-06-02.md` , `ai-smell-2026-06-02.md` , `SKILL-USAGE-*.json`); later same-skill arms occasionally reused them (flagged on **ai-codegen** Haiku-B). These stray files should be cleaned up.
- **Auto-synthesized tasks + single-agent orchestrator approximation** inflate `audit` / `readme` skill-arm cost.
- **Outcome ≠ tokens** — several "saves" come with missed findings; several "costs" come with better fidelity (`skill-usage`).
- **Total measurement cost:** ~2.54M `subagent_tokens` across 48 arms (Opus ~931K, Sonnet ~856K, Haiku ~751K).
- Cross-links: first calibration → `docs/audits/skill-calibration-2026-06-01.md` ; optimization study → `docs/audits/skills-optim-study-2026-05-31.md` .